

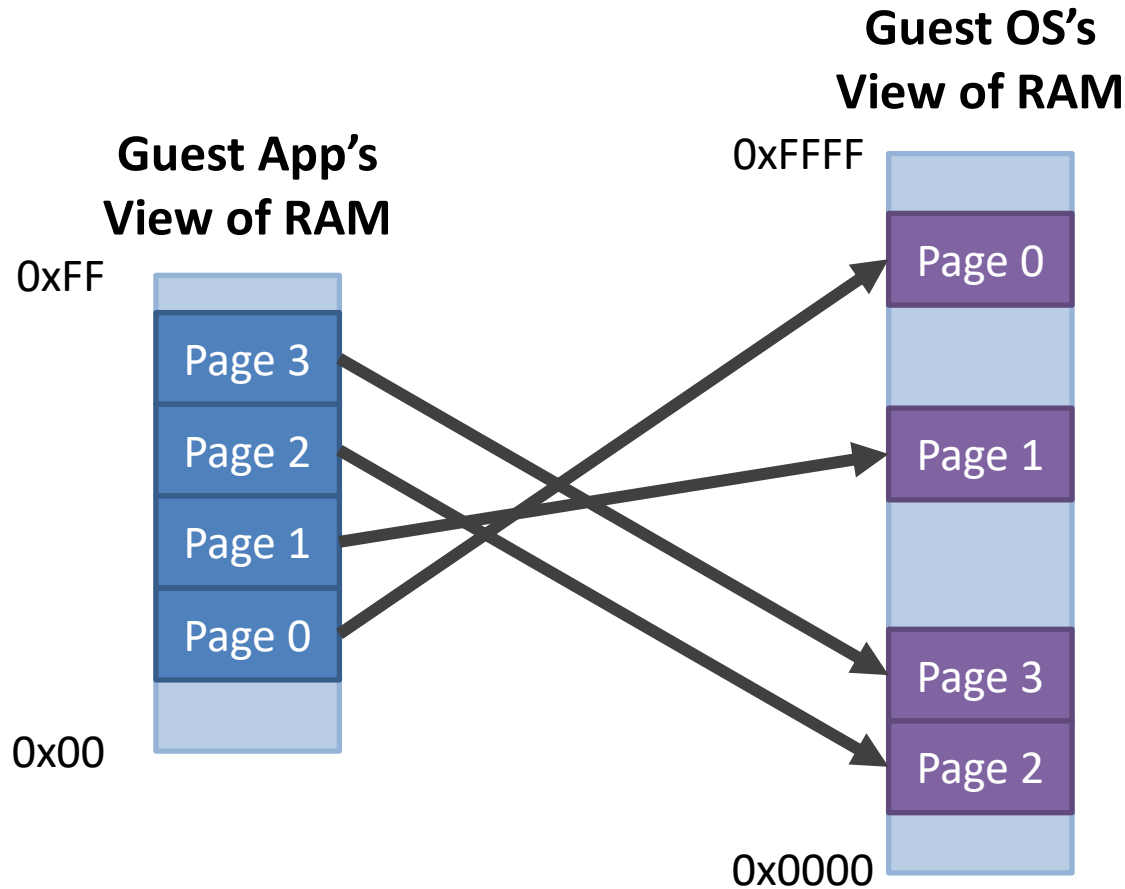
x86 Memory Virtualization

Dr. Amit Praseed

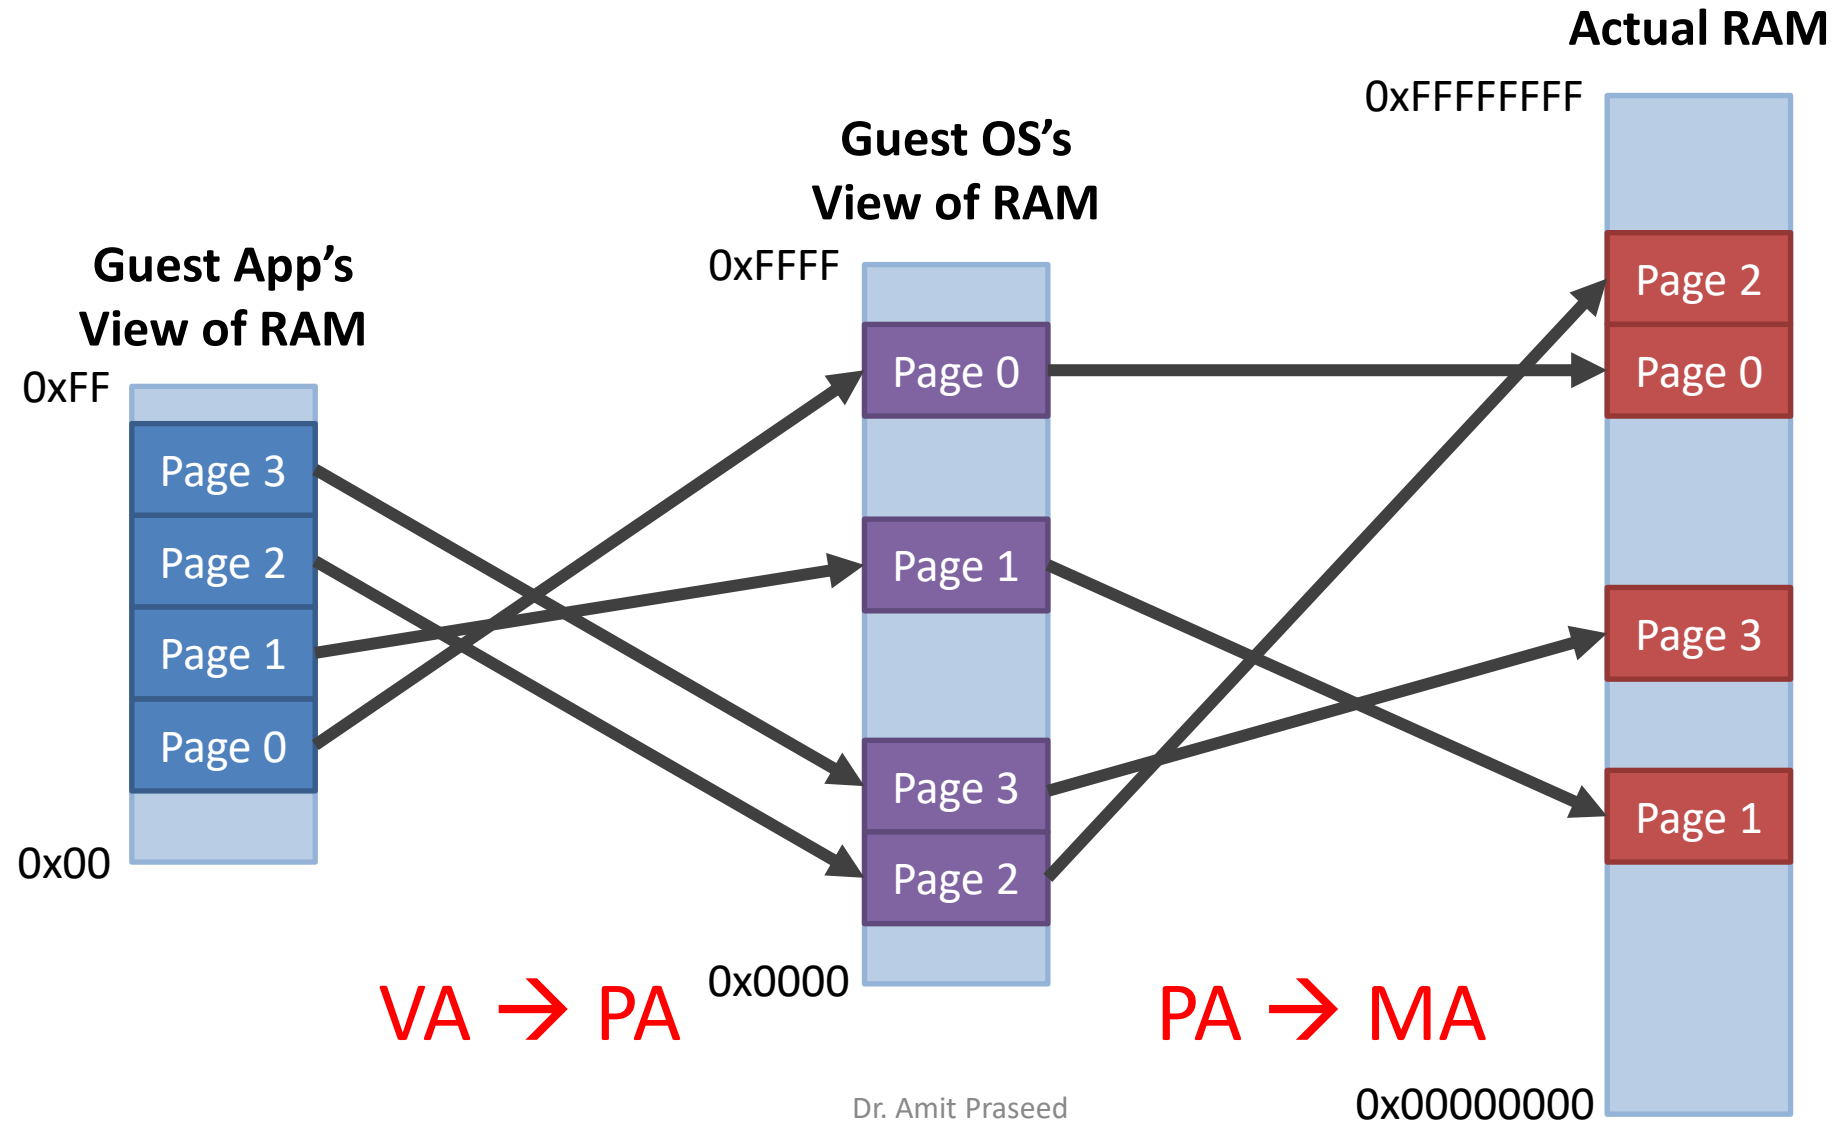
Memory Virtualization

- Typical x86 architecture has a virtual to physical address mapping ($VA \rightarrow PA$)
- Virtualized x86 architecture requires a two level address translation
 - $VA \rightarrow PA$
 - Physical address (PA) \rightarrow Machine Address (MA)
- Guest OS has no idea about this translation
 - Guest continues to maintain page tables containing $VA \rightarrow PA$ mappings

Paging without Virtualization



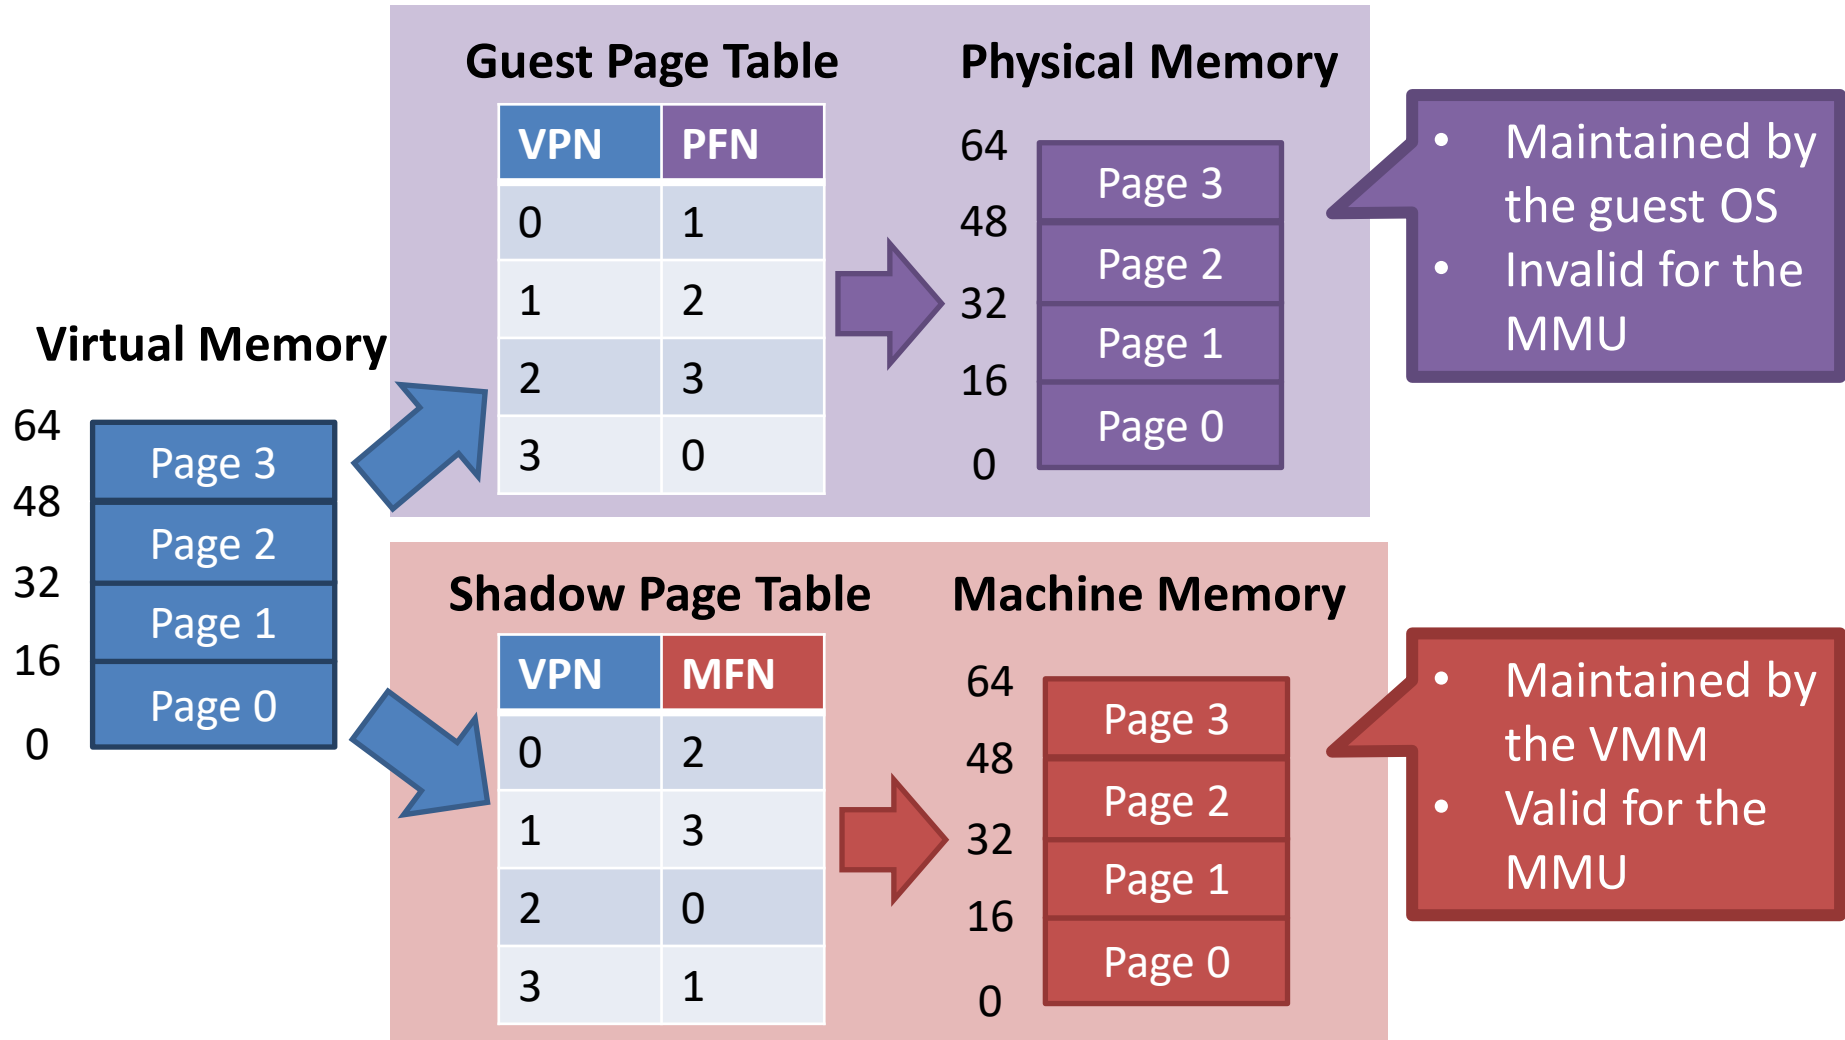
Paging with Virtualization



Does a 2 Level Indirection Work?

- Guest is only aware of VA \rightarrow PA mapping
 - Issues only VA to hardware MMU
 - **MMU supports only a single mapping**
- Solution: Shadow Page Table
 - Hypervisor maintains a single shadow page table in the MMU
 - Shadow page table contains direct VA \rightarrow MA mapping
 - Trick is to maintain consistency!!!

Shadow Page Tables



Building Shadow Page Tables

- The guest can update its page tables at any time
 - Not a privileged instruction – not trapped!
 - Without knowing when the guest OS updates its page table, the hypervisor cannot maintain the correct entry in the shadow page table
- Solution: Mark the guest page tables as read only
 - Writing generates an exception, which can be trapped by hypervisor

What happens during page faults?

- Two kinds of page faults can occur:
 - True Miss : The mapping does not exist in the guest page table
 - Hidden Miss : The mapping exists in the guest page table, but is absent in the shadow page table
- The hypervisor should disambiguate between the two
- On every miss, the hypervisor walks the guest page table [Tracing]
 - If a mapping exists, the hypervisor silently updates the shadow page table and retries the instruction
 - Otherwise, the hypervisor forwards the page fault to the guest OS for handling

Pros and Cons

- The good: shadow tables allow the MMU to directly translate guest VPNs to hardware pages
 - Thus, guest OS code and guest apps can execute directly on the CPU
- The bad:
 - Double the amount of memory used for page tables
 - i.e. the guest's tables and the shadow tables
 - Overhead due to VMM traps
 - TLB flush whenever a “world switch” occurs
 - Loss of performance

Second Level Address Translation (SLAT)

- Hardware support for memory virtualization
 - Extended Page Tables (EPT) – Intel
 - Nested Page Tables (NPT) - AMD
- Walking the guest and host page tables can be combined into a single multilevel page table
 - Page table depth increases tremendously in some cases!!!
 - Extremely important to use an effective TLB to avoid expensive page table walks
- TLB modified to reduce miss rate
 - Larger TLB
 - More expensive, though!!!
 - Tagged TLB
 - Every entry in TLB now has an address space identifier
 - No need to flush TLB for every “world switch”