

Big Data and its Implications on the Cloud


Dr. Amit Praseed

How much data do we generate?

- **Structured Data**
 - Relational Databases
 - Well defined schema
- **Unstructured Data**
 - Videos, audio, images etc.
- **Semi-structured Data**
 - structured in form but not well defined (no schema)
 - XML files



The Rise of Unstructured Data



THE 3Vs OF BIG DATA

VOLUME

- ◆ Amount of data generated
- ◆ Online & offline transactions
- ◆ In kilobytes or terabytes
- ◆ Saved in records, tables, files



VELOCITY

- ◆ Speed of generating data
- ◆ Generated in real-time
- ◆ Online and offline data
- ◆ In Streams, batch or bits



VARIETY

- ◆ Structured & unstructured
- ◆ Online images & videos
- ◆ Human generated - texts
- ◆ Machine generated - readings



The Rise of NoSQL

- NoSQL = Not only SQL
- A fundamental shift or alternative to storing data which does not conform to the relational format
- Features
 - Schema Agnostic
 - Non relational
 - Commodity hardware
 - Highly distributable

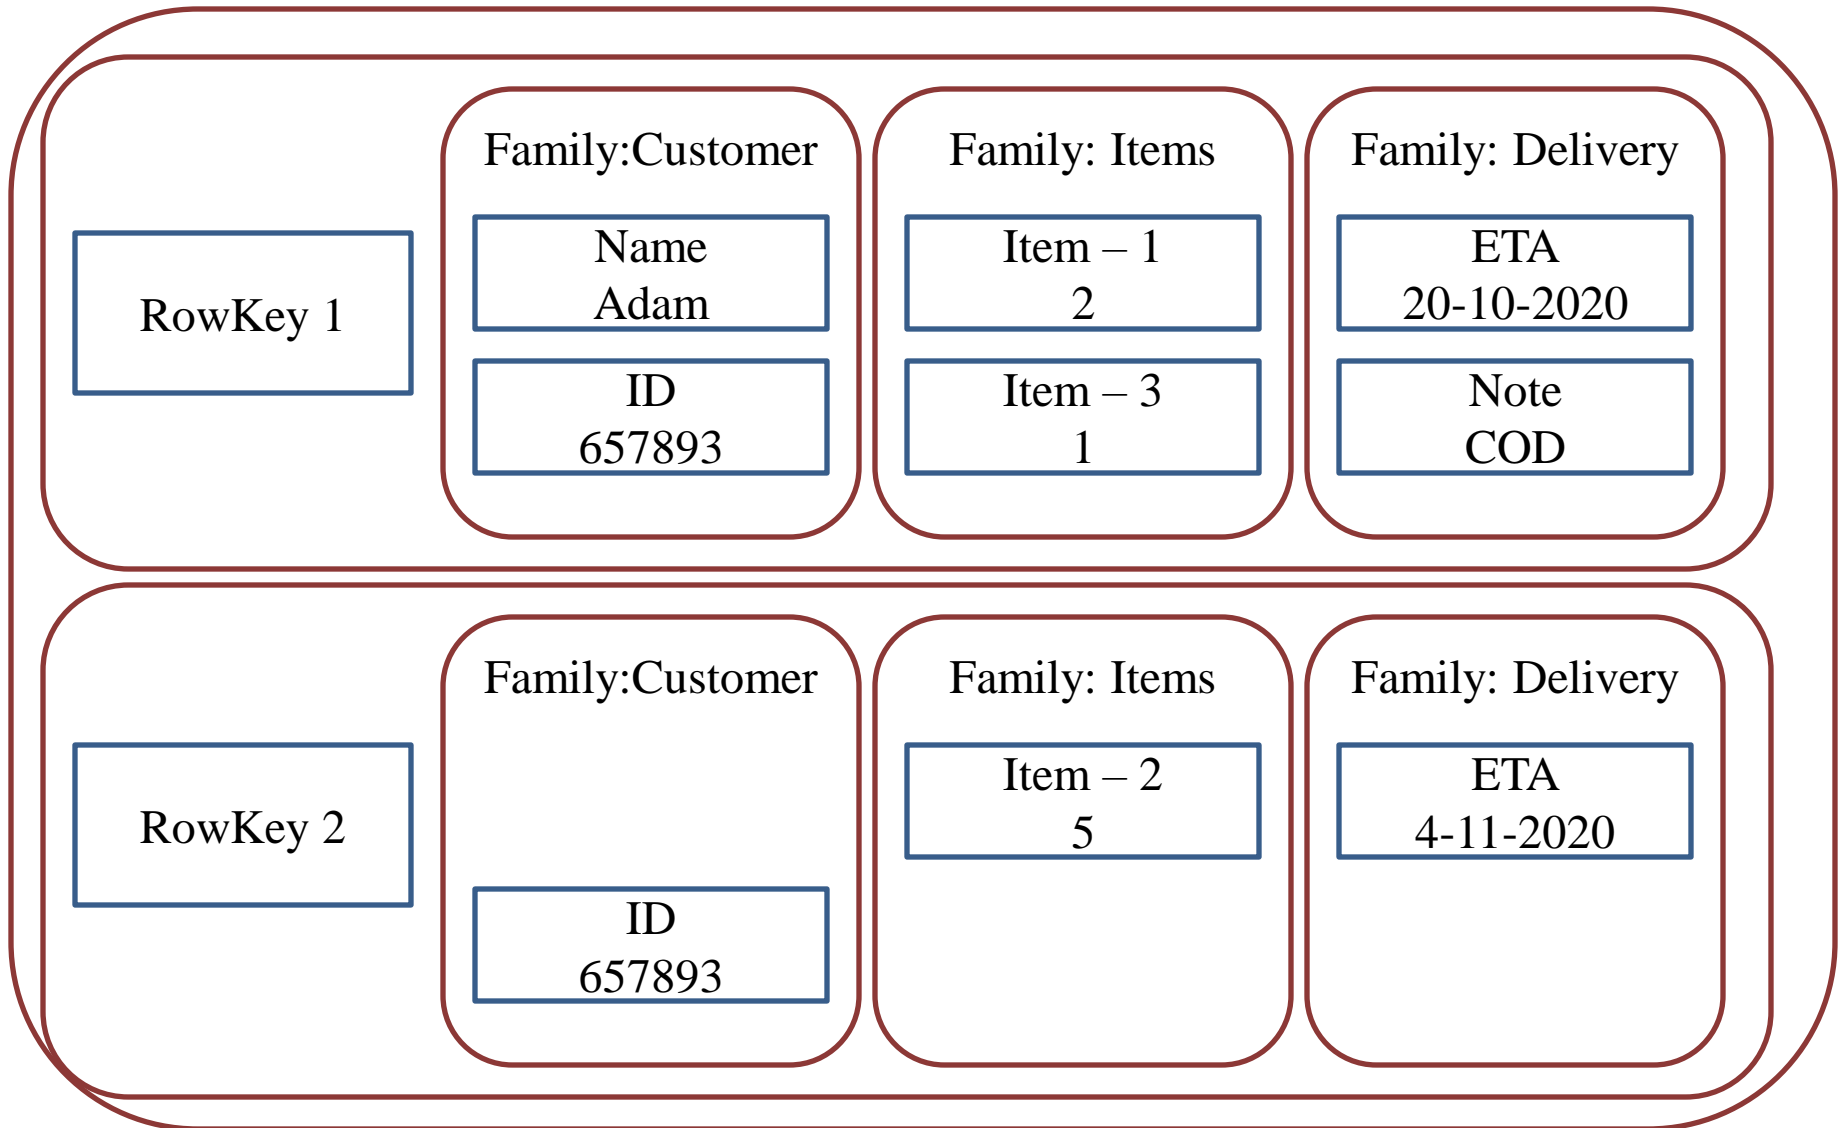
Four Core NoSQL Varieties

- Columnar
- Key-Value
- Triple
- Document

Columnar Databases

- Similar to relational model – concept of rows and columns still exist
- Optimized and designed for faster column access
- Ideal for running aggregate functions or for looking up records that match multiple columns
- A single record may consist of an ID field and multiple column families
- Each one of these column families consists of several fields. One of these column families may have multiple “rows”

Columnar Databases




Key – Value Stores

- An ID field — the key in key-value stores — and a set of data
- Some key-value stores support typing (such as integers, strings, and Booleans) and more complex structures for values (such as maps and lists)
- Key-value stores are optimized for speed of ingestion and retrieval

Triple and Graph Stores

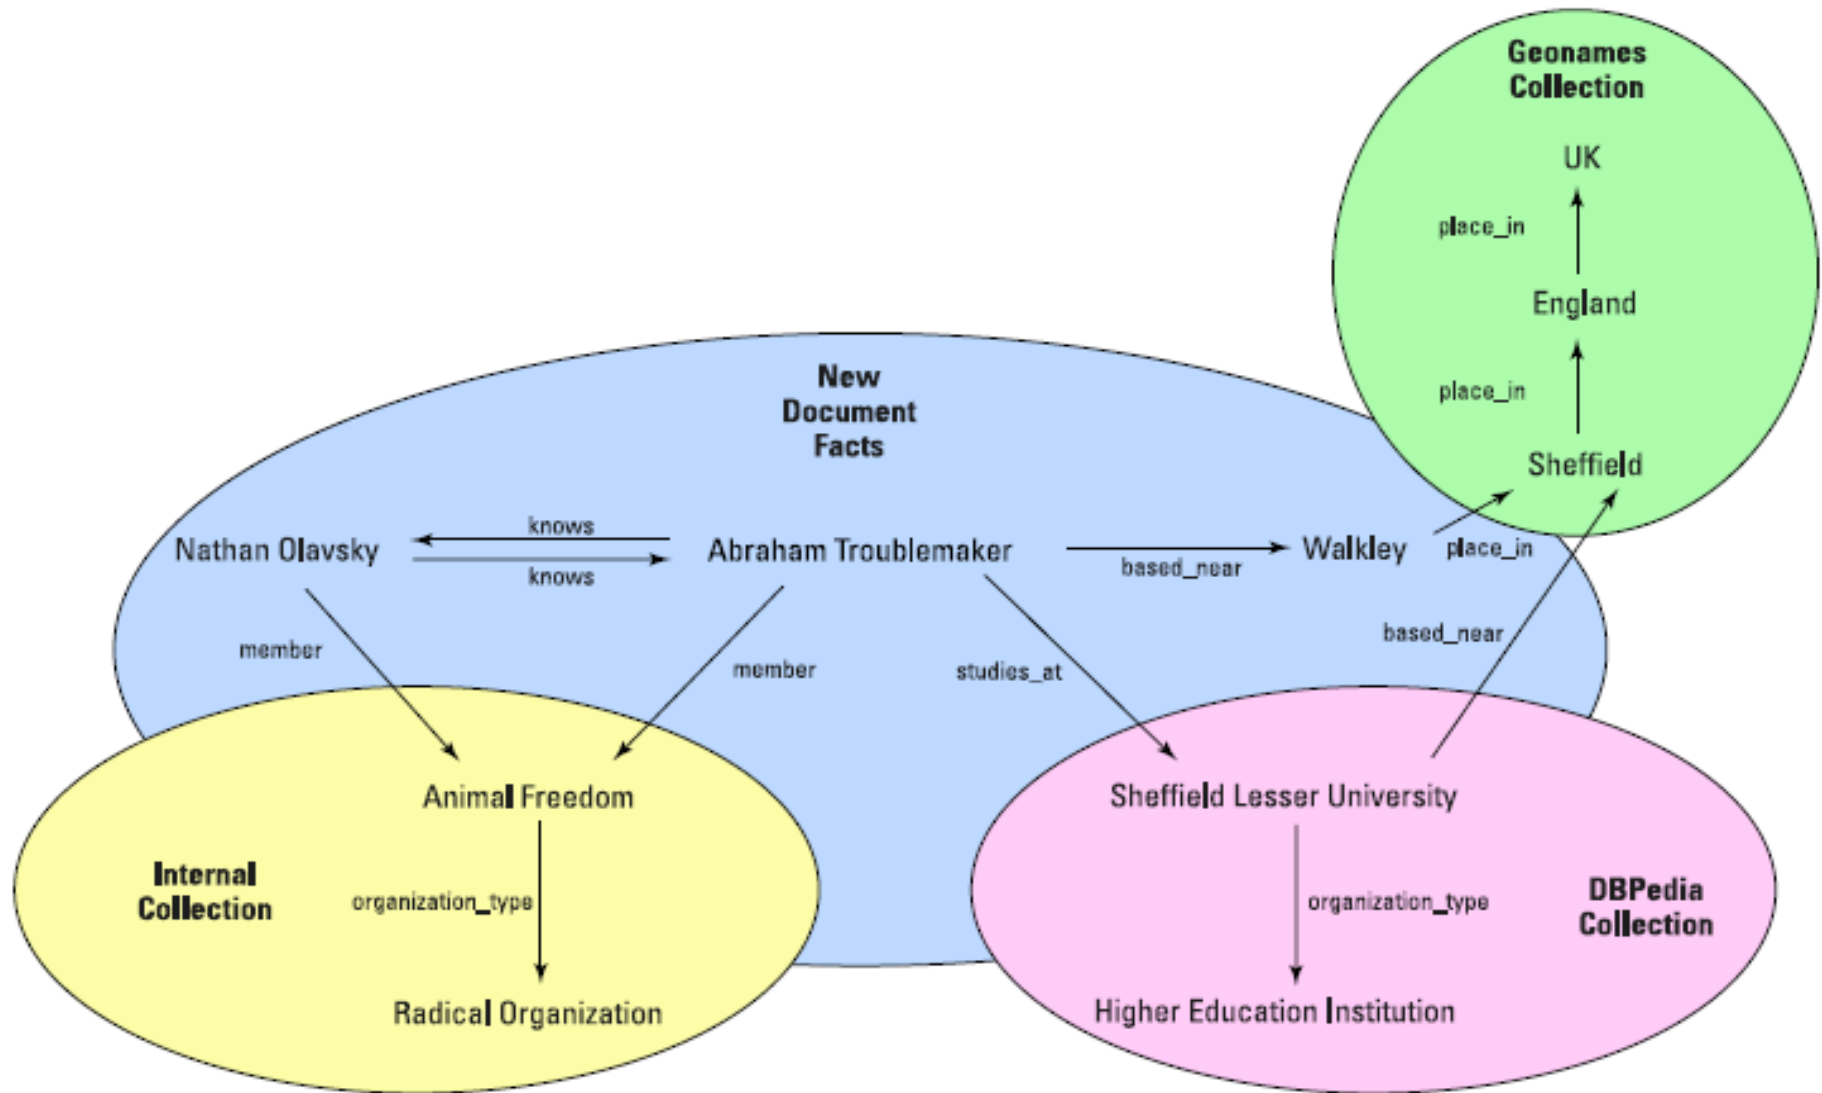
- Every *fact* or assertion is described as a triple of subject, predicate, and object:
 - A *subject* is the thing you're describing. It has a unique ID called an IRI. It may also have a type, which could be a physical object (like a person) or a concept (like a meeting).
 - A *predicate* is the property or relationship belonging to the subject. This again is a unique IRI that is used for all subjects with this property.
 - An *object* is the intrinsic value of a property (such as integer or Boolean, text) or another subject IRI for the target of a relationship.

Triple and Graph Stores



- Three points of information
 - Adam is a person
 - Cheese is a food item
 - Adam likes cheese

Triple and Graph Stores



Document Stores

- Hold documents that combine information in a single logical unit
- Retrieving all information from a single document is easier with a database and is more logical for applications
- A document is any unstructured or tree-structured piece of information.
 - It could be a recipe, financial services trade, PowerPoint file, PDF, plain text, or JSON or XML document.

Examples of NoSQL Products

- **Columnar:** DataStax, Apache Cassandra, HBase, Apache Accumulo, Hypertable
- **Key-value:** Basho Riak, Redis, Voldemort, Aerospike, Oracle NoSQL
- **Triple/graph:** Neo4j, Ontotext's GraphDB (formerly OWLIM), MarkLogic, OrientDB, AllegroGraph, YarcData
- **Document:** MongoDB, MarkLogic, CouchDB, FoundationDB, IBM Cloudant, Couchbase

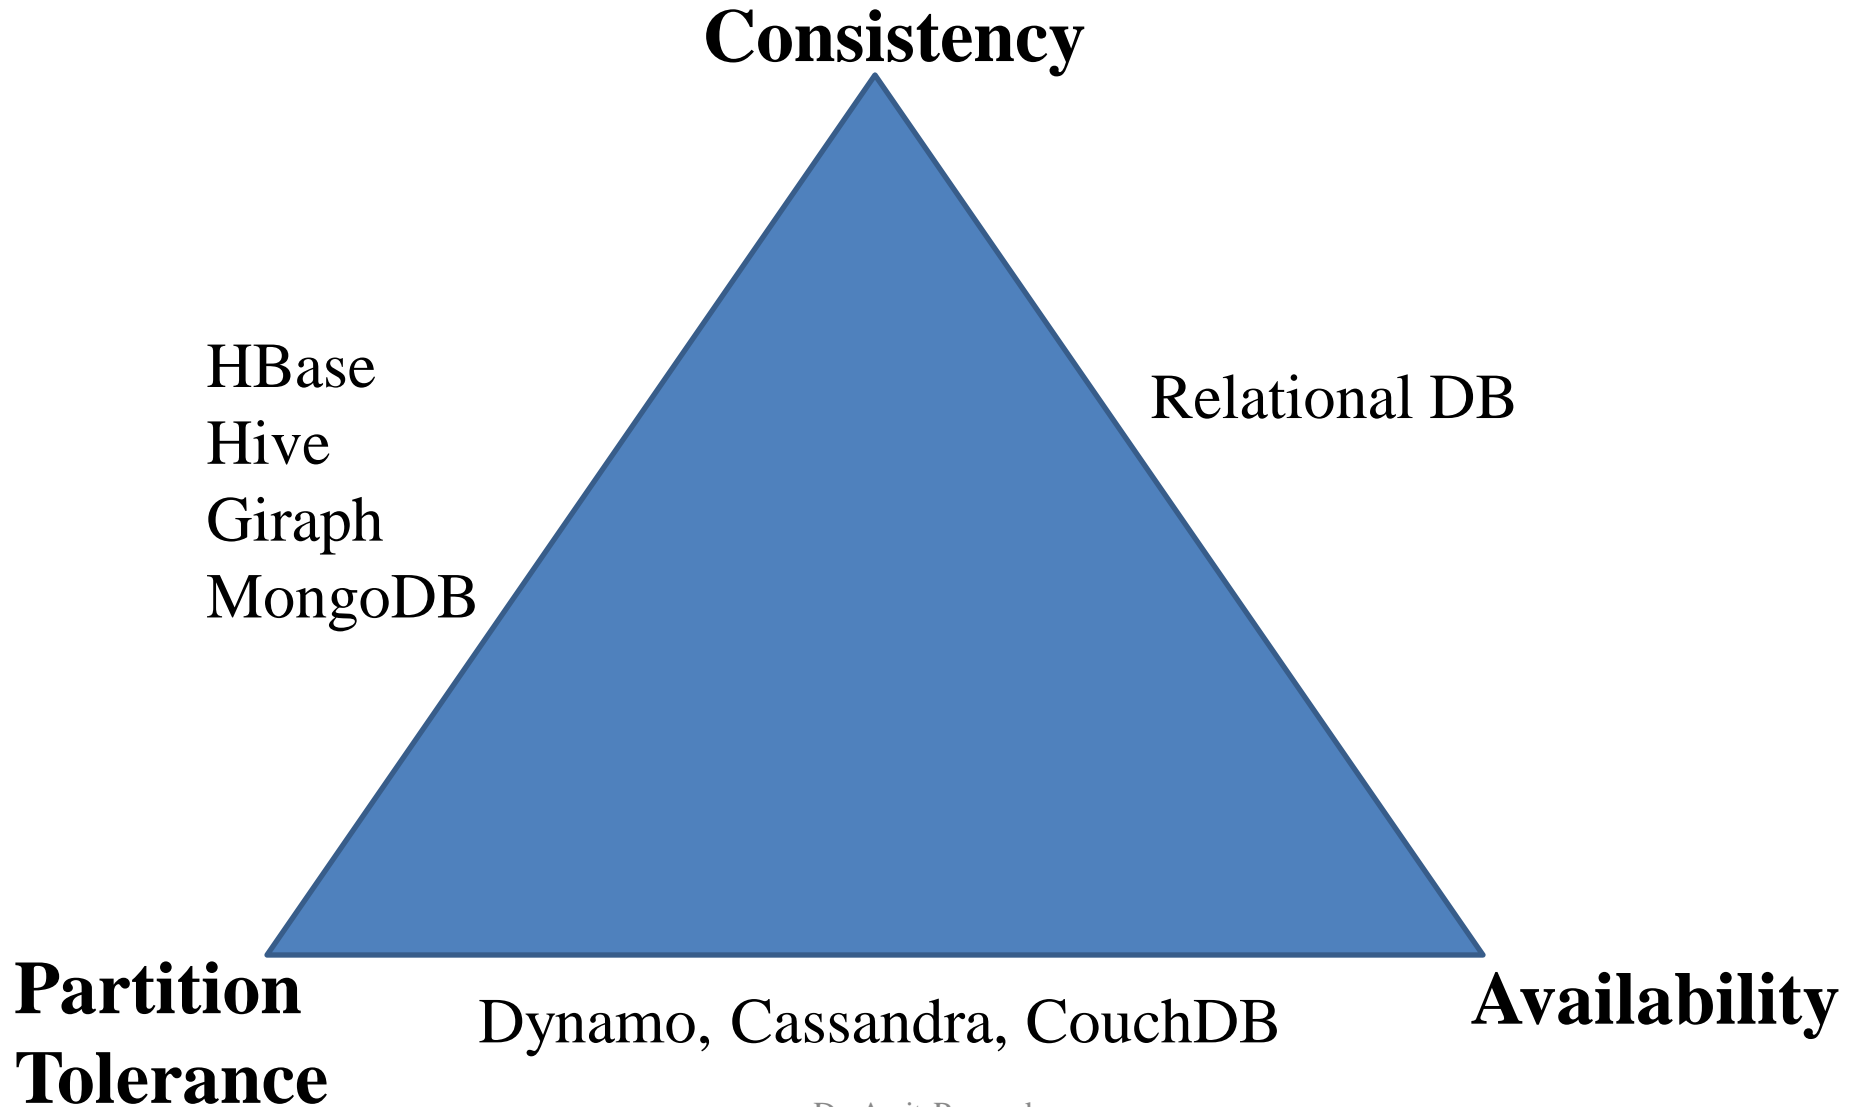
The CAP Theorem

- It is impossible for a distributed data store to simultaneously provide more than two out of the following three guarantees:
 - **Consistency:** Every read receives the most recent write or an error
 - **Availability:** Every request receives a response, without the guarantee that it contains the most recent write
 - **Partition tolerance:** The system continues to operate despite an arbitrary number of messages being dropped by the network between nodes

Consistency in NoSQL

- NoSQL relaxes the ACID consistency model used in relational databases
- NoSQL uses an eventual consistency model called BASE
 - Basically Available
 - Soft state
 - Eventual consistency

ACID, BASE and the CAP Theorem



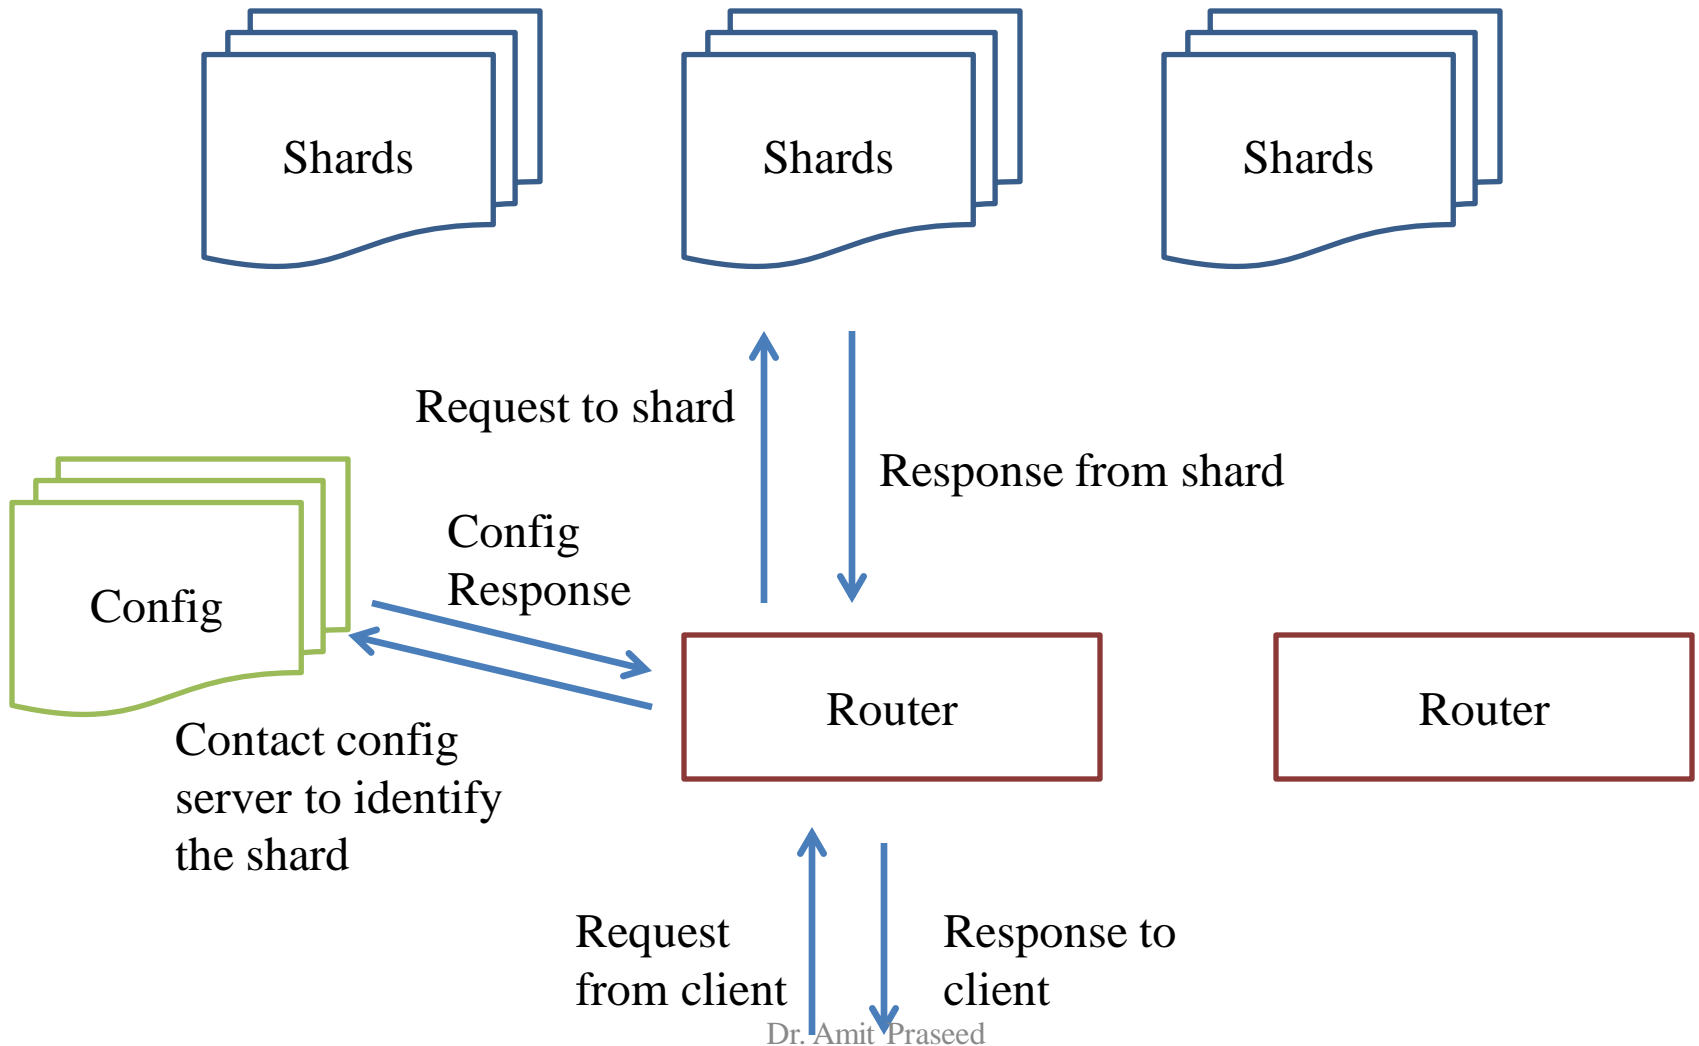
MongoDB

- A document-oriented, NoSQL database
 - Hash-based, schema-less database
 - Atomic writes and fully-consistent reads
 - Master-slave replication with automated failover (replica sets)
 - Built-in horizontal scaling via automated range-based partitioning of data (sharding)
 - No joins nor transactions
 - Consistency + Partition Tolerance

MongoDB Terminology

- A MongoDB instance may have zero or more ‘databases’
- A database may have zero or more ‘collections’
- A collection may have zero or more ‘documents’
- A document may have one or more ‘fields’
- MongoDB ‘Indexes’ function much like their RDBMS counterparts.

MongoDB Architecture



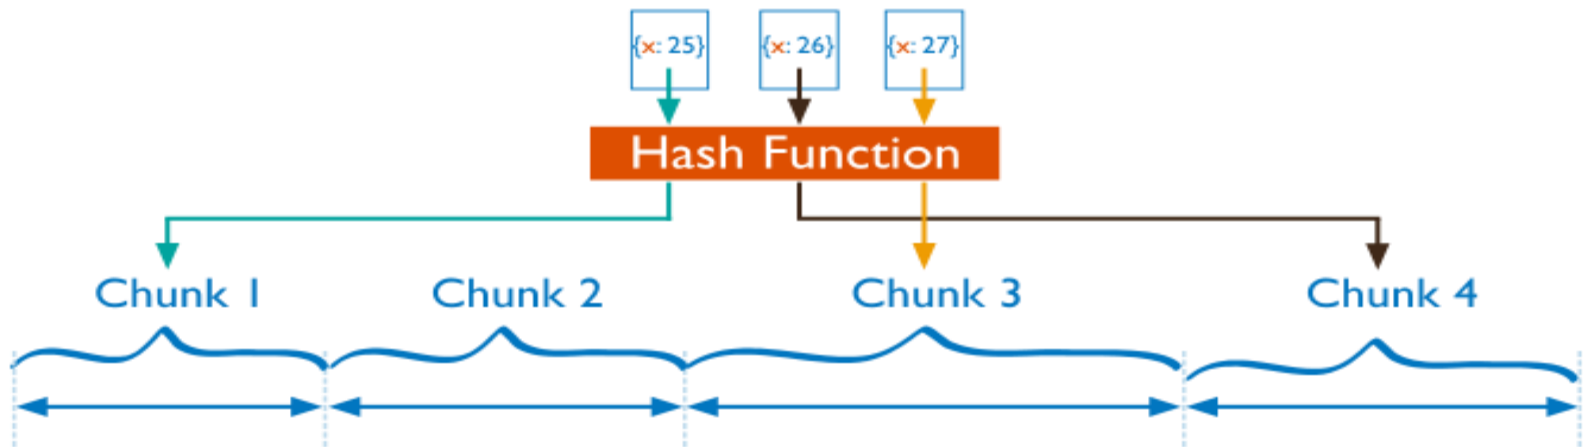
MongoDB Replication

- A replica set is a group of mongod instances that maintain the same data set.
 - A replica set contains several data bearing nodes and optionally one arbiter node.
 - Of the data bearing nodes, one member is deemed the primary node, while the other nodes are deemed secondary nodes.
- The primary node receives all write operations.
 - The primary records all changes to its data sets in its operation log, i.e. oplog
 - The secondaries replicate the primary's oplog and apply the operations to their data sets

Sharding in MongoDB

- **Hashed Sharding**

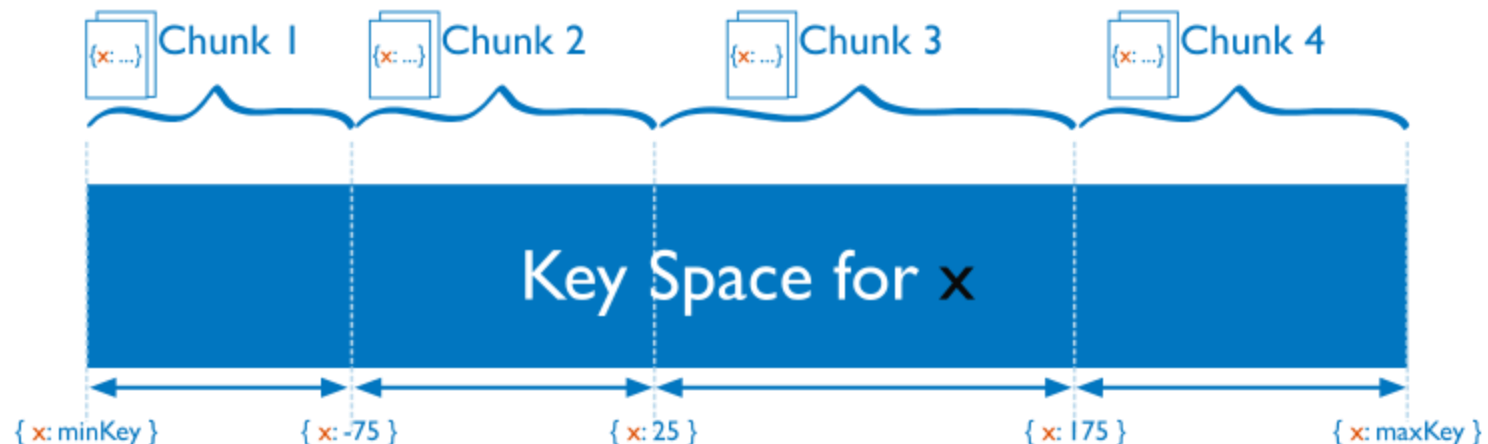
- Involves computing a hash of the shard key field's value.
- Each chunk is then assigned a range based on the hashed shard key values.
- Even data distribution, but may result in more cluster wide broadcast operations



Sharding in MongoDB

- **Ranged Sharding**

- Involves dividing data into ranges based on the shard key values.
- Each chunk is then assigned a range based on the shard key values
- Allow targeted operations, but poorly chosen shard key may result in uneven data distribution

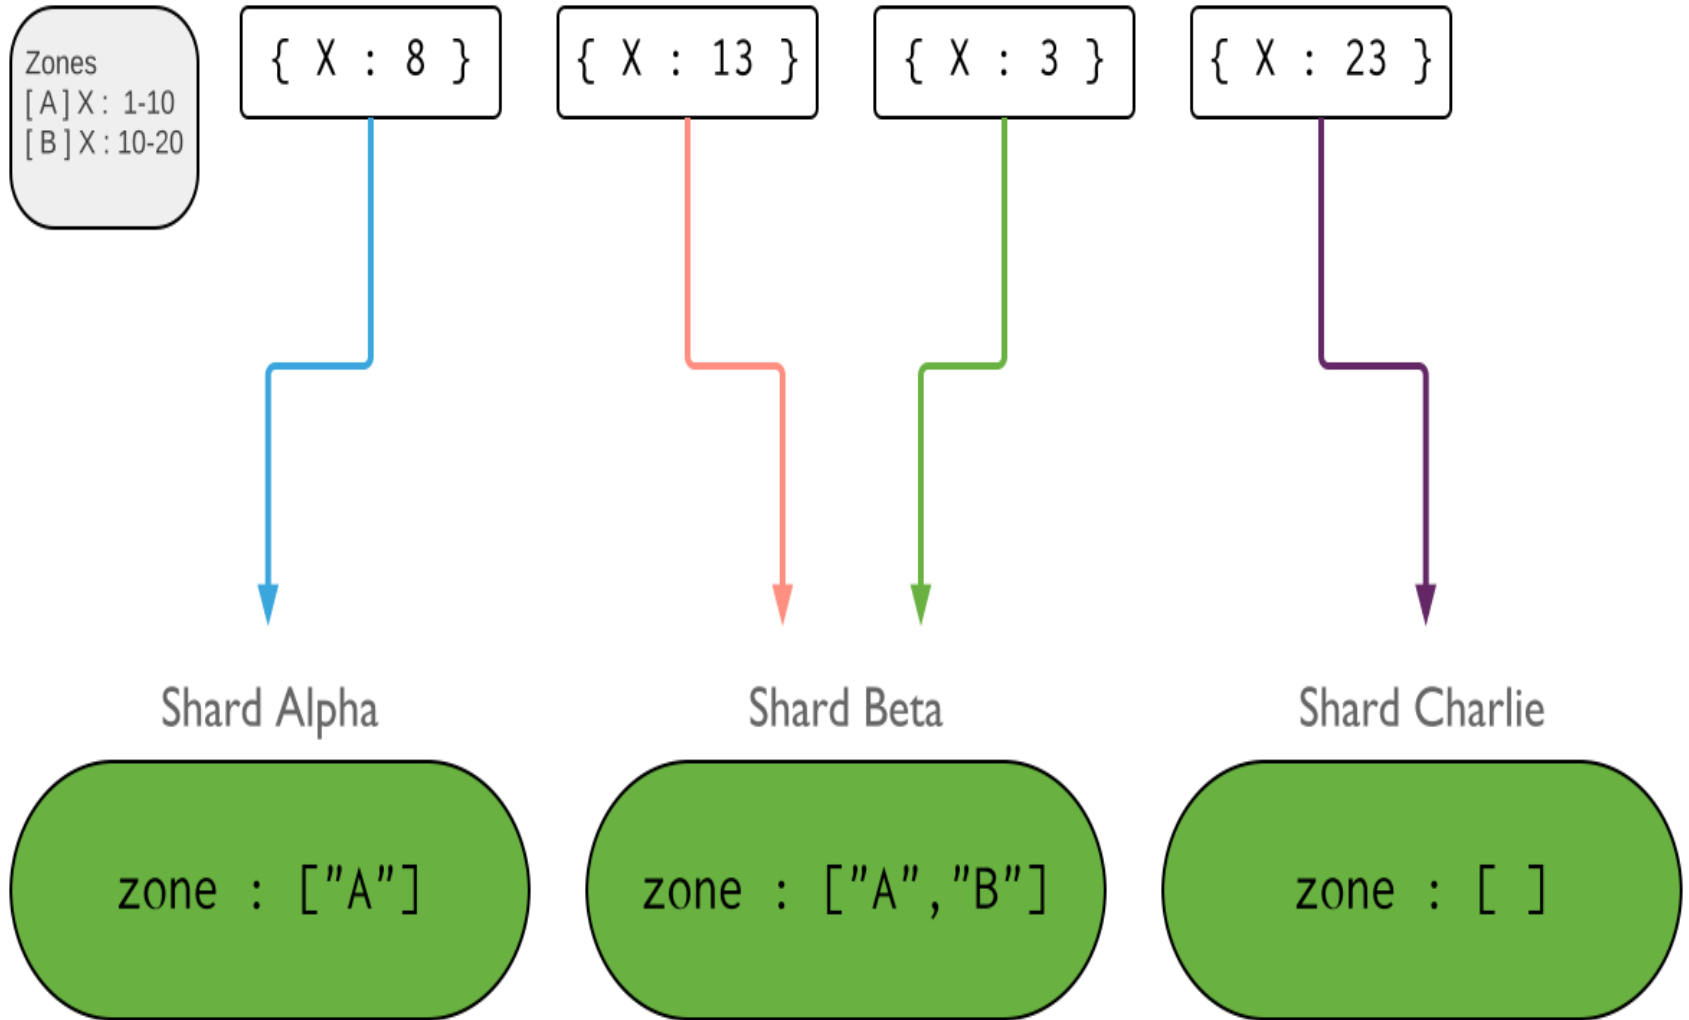


Sharding in MongoDB

- **Zoned Sharding**

- Zones can help improve the locality of data for sharded clusters that span multiple data centers.
- In sharded clusters, you can create zones of sharded data based on the shard key.
- You can associate each zone with one or more shards in the cluster. A shard can associate with any number of zones.
- In a balanced cluster, MongoDB migrates chunks covered by a zone only to those shards associated with the zone.
- Each zone covers one or more ranges of shard key values

Zoned Sharding



Choosing a Shard Key

- The shard key is either an indexed field or indexed compound fields that determines the distribution of the collection's documents among the cluster's shards.
- All sharded collections must have an index that supports the shard key; i.e. the index can be an index on the shard key or a compound index where the shard key is a prefix of the index
- The choice of shard key affects the creation and distribution of the chunks across the available shards. This affects the overall efficiency and performance of operations within the sharded cluster.